

# purrr solutions

## *Jumping Rivers*

First we must load the **tidyverse**

```
library("tidyverse")
```

### Question 1

```
data(1, package = "jrTidyverse2")
```

a) 1 is a list with 3 vector elements; A, B, C. Each vector contains the heights of a different class at a university.

- How many students are in each class?
- Which class has the highest average height?
- Which class has the smallest student?
- Which class has the tallest student?

Hint: use `map()`, `map_dbl()`, or `map_df()`

```
map_dbl(1, length)
## A B C
## 28 30 24
map_dbl(1, mean)
## A B C
## 176.4357 168.9733 181.8833
map_dbl(1, min)
## A B C
## 154.1 159.1 154.7
map_dbl(1, max)
## A B C
## 199.8 176.8 215.6

### OR

map_df(1, ~tibble(length = length(.x),
                  mean = mean(.x),
                  min = min(.x),
                  max = max(.x)))
## # A tibble: 3 x 4
## length mean min max
## <int> <dbl> <dbl> <dbl>
## 1 28 176. 154. 200.
## 2 30 169. 159. 177.
## 3 24 182. 155. 216.
```

b) For each class, work out the lower and upper confidence interval bounds for the mean given that they are  $\mu - 1.96\sigma$  and  $\mu + 1.96\sigma$  respectively. Where  $\mu$  is the sample mean and  $\sigma$  is the sample standard deviation.

Hint: use the formula notation and either `map()`, `map_dbl()` or `map_df()`.

```
map_df(1, ~tibble(lower = mean(.x) - 1.96*sd(.x),
                    upper = mean(.x) + 1.96*sd(.x)))
## # A tibble: 3 x 2
##   lower upper
##   <dbl> <dbl>
## 1  155.  198.
## 2  160.  178.
## 3  144.  220.
```

## Question 2

Now we're going to look at a list containing happiness rankings for countries around the globe.

```
data(happiness, package = "jrTidyverse2")
```

- a) How long is the list? Is this a recursive list? How many countries does the list contain information on? For each country how many pieces of information is there?  
Hint: use `str()`

```
str(happiness, max.level = 0)
## List of 146
# 146 element in the list
str(happiness, max.level = 1, list.len = 3)
## List of 146
## $ :List of 12
## $ :List of 12
## $ :List of 12
## [list output truncated]
# Yes, recursive list
str(happiness, max.level = 2, list.len = 3)
## List of 146
## $ :List of 12
## ..$ Country           : chr "Switzerland"
## ..$ Region            : chr "Western Europe"
## ..$ Year               : num [1:3] 2015 2016 2017
## .. [list output truncated]
## $ :List of 12
## ..$ Country           : chr "Iceland"
## ..$ Region            : chr "Western Europe"
## ..$ Year               : num [1:3] 2015 2016 2017
## .. [list output truncated]
## $ :List of 12
## ..$ Country           : chr "Denmark"
## ..$ Region            : chr "Western Europe"
## ..$ Year               : num [1:3] 2015 2016 2017
## .. [list output truncated]
## [list output truncated]
# Each element of the list is another list representing a country.
# Each list contains elements representative of happiness information
# on that country for three successive years. Therefore there is
# 146 countries and 12 vectors of information on each.
```

- b) Return the name of each country contained in the list. To make it a bit easier to read return the output as a character vector.

```
country_names = map_chr(happiness, "Country")
```

- c) Try `names(happiness)`, what happens? Use the answer to b) to rename each element of the list after it's representative country.

```
names(happiness) = country_names
```

- d) What has the UK's average happiness rank been over the last 3 years?

```
UK_rank = happiness[["United Kingdom"]]$`Happiness Rank`
mean(UK_rank)
## [1] 21
```

- e) Over the last 3 years, what is the average happiness score for every country? Store this in a vector of doubles.

```
mean_hap = happiness %>%
  map_dbl(~ mean(.x[["Happiness Score"]]))
head(mean_hap)
## Switzerland      Iceland      Denmark      Norway      Canada      Finland
##      7.530000      7.522000      7.525000      7.519000      7.382333      7.429333
```

- f) Which region of the world has the high average happiness score?  
Hint: store the region for each country in a vector, combine it into a data frame with the average happiness then use `dplyr`.

```
region = happiness %>%
  map_chr("Region")

region_hap = data.frame(region = region, mean_hap = mean_hap)

region_hap %>%
  group_by(region) %>%
  summarise(av_region_hap = mean(mean_hap)) %>%
  arrange(av_region_hap)
## # A tibble: 10 x 2
##   region                av_region_hap
##   <fct>                <dbl>
## 1 Sub-Saharan Africa      4.11
## 2 Southern Asia           4.59
## 3 Middle East and Northern Africa  5.36
## 4 Central and Eastern Europe  5.37
## 5 Southeastern Asia       5.40
## 6 Eastern Asia            5.49
## 7 Latin America and Caribbean  6.05
## 8 Western Europe          6.69
## 9 North America           7.23
## 10 Australia and New Zealand  7.30
```

- g) Using `ggplot2` and `geom_col()`, plot the answer to f) as a bar chart

```
region_hap %>%
  group_by(region) %>%
  summarise(av_region_hap = mean(mean_hap)) %>%
  arrange(av_region_hap) %>%
  ggplot() +
  geom_col(aes(x = region, y = av_region_hap)) +
```

```
coord_flip()
```

### Question 3

Load the data and the required packages

```
library("broom")
data(beer_tidy, package = "jrTidyverse2")
```

Here we have a data set of around 500 beers along with their alcohol percentage, colour and type.

```
head(beer_tidy)
## # A tibble: 6 x 4
##   URL                                ABV Color Type
##   <chr>                            <dbl> <dbl> <chr>
## 1 Deschutes Mirror Pond Pale Ale Clone  5.08  6.58 Ale
## 2 Kiwanda Cream Ale                    5.21  3.08 Ale
## 3 Christmas Ale                       6.82 24.0  Ale
## 4 Strong Scotch Ale Smuttynose Clone   8.96 21.1  Ale
## 5 Creamy Cream Ale                     4.96  3.07 Ale
## 6 Lemon Ale                           5.31  9.14 Ale
```

We're going to run a linear regression on the data testing how colour affects alcohol percentage. However, using **purrr**, we are going to do it for each different type of beer.

- 1) Nest the beers within each type. Save this data in a variable called **beer\_nest**.

```
beer_nest = beer_tidy %>%
  nest(-Type)
```

- 2) For each type of beer, fit a linear regression model with alcohol percentage as the response variable and colour as the predictor. Store the models in a list column next to your nested data. To help you out, the below code is how you would fit the model to the entire data set, before nesting.

```
fit = lm(ABV ~ Color, data = beer_tidy)
```

Hint: use **mutate()** and **map()**

```
beer_nest = beer_nest %>%
  mutate(fit = map(data, ~lm(ABV ~ Color, data = .x)))
```

- 3) For each model, mutate a new column with the tidied models in. We want to plot the points eventually, so we want to tidy the model such that it returns the fitted points.

Hint: use **mutate()**, **map()** and **augment()**

```
beer_nest = beer_nest %>%
  mutate(tidyfit = map(fit, ~augment(.x)))
```

- 4) Select the columns containing the beer type and the tidied models and unnest the data.

```
beer_nest = beer_nest %>%
  select(Type, tidyfit) %>%
  unnest()
```

- 5) Given you have stored unnested models in the variable **beer\_nest()**, the following code WILL plot the lines made by the linear regression models

```
ggplot(beer_nest) +
  geom_line(aes(x = Color, y = .fitted, colour = Type), size = 2)
```